

Lipsync starring yourself

The user picks a template, uploads an ordinary selfie, picks a style — and gets a lipsync with themselves on screen. Six finished templates below: client photo on the left, template aesthetic in the middle, a frame of the result on the right.

What we are offering

Two things, not one. The second matters more: it turns a one-off effect into a family that can be extended without us.

FIRST

The generation pipeline

Eight stages from selfie to a finished vertical clip with sound. Seven are free and sit before the paid one, so defects are caught on a still image for a fraction of a cent rather than on a dollar of video. Every check has three outcomes and prints numbers next to the verdict.

SECOND

The template-building instrument

The author writes a prompt, supplies the demo identity — and gets an aesthetic ready to ship. The framing is measured off the driving video and reaches the prompt on its own. A new template takes one run instead of a sprint, and needs no engineer standing by.



The family: six templates, one machine

Each triptych reads left to right. The first two images are inputs, the third is what the user received. There are two clients: one male selfie and one female, both generated to look like ordinary phone photos.

CLIENT PHOTO



TEMPLATE AESTHETIC

RESULT



Y2K

f • driving b4 • identity 0.9504




Mid-century

f • driving b4 • identity 0.8761



Ice cream

f • driving b5 • identity 0.8188



Country

f • driving b5 • identity 0.8715



Fisheye

m • driving b2 • identity 0.6287



Tomatoes

m • driving b3 • identity 0.6648

What is identical across all six

This is the machine. Six different aesthetics, six different dances — one predictable output. Not one technical axis failed even once.

720×1280

exactly 9:16 out of
the model, no
cropping

300

frames, exactly 10.0
seconds in all six

0 %

of frame area lost in
assembly

0

editing cuts in the
frame

\$0.70

per clip, measured
four times

Where quality dips, and why

On some clips the face drifts a little. We see it and say so plainly, because the cause is measured and is fixed by choosing better material, not by rebuilding the pipeline.

It comes down to how large the face is in the driving video. Where the source face is large, identity holds. Where it is three times smaller, ArcFace stops being able to judge it and the model stops carrying it confidently.

DRIVING	FACE IN FRAME	FRAMES BELOW THE BAR	IDENTITY ON OUTPUT
b3 ♂	96–139 px	13 %	0.6648 eye judges
b2 ♂	85–110 px	87 %	0.6287 eye judges
b4 ♀	34–56 px	80 % + 20 % no face	0.876 / 0.950 drifts
b5 ♀	34–53 px	94 % + 6 % no face	0.819 / 0.872 drifts

The match is exact: both drivings with a small face produced drifting identity, both with a large face held it. Two cheap fixes: **more careful driving selection** or a **manual crop** to the face size needed.

Requirements for the driving video

The driving video is the only input we do not draw but buy. So its requirements sit separately: this is a purchasing checklist, not a setting.

Face larger than 100 px

Below that there is nothing to measure identity with, and the model carries it unconfidently. Checked automatically, before generation.

One scene, zero editing cuts

A cut inside the window becomes an instant shot change in the finished clip. Checked automatically on brightness, before poses are read.

Scene longer than the product length

The window is cut from the middle of the longest scene with equal margins on both sides.

The person stays inside the frame

The driving framing is carried into the aesthetic prompt: the model maps the pose from the video, so if the person leaves the frame there, they leave it in the result too.

Lipsyncs are the hardest material

Of all formats lipsync is the hardest, and hardest of all when the character **turns around**. On a turn the face leaves the frame entirely: there is nothing left to measure identity on, and nothing for the model to hold. Both of our female drivings contain turns, and that is the second half of the explanation for their numbers.

What Kling limits, and what to do about it

The two ceilings we hit are properties of one model, not of the approach. Both lift by swapping the engine, with no change to the logic.

MODEL LIMITATION

Three-second minimum. Anything shorter is simply refused, in the model's own wording in the API response.

Unstable frame-level output. Duration and frame accuracy have to be verified on the output rather than set on the input.

WHAT INSTEAD

Any other video model over an API – the pipeline logic does not depend on the engine; one call changes.

Or open source on our own GPUs: a parallel pipeline is already built for that route, where duration and frames are ours to set, not the vendor's.

The expensive Kling tier was ruled out by a number: **\$0.90** per second against **\$0.07 – 12.8×** the cost, and it won on no axis at all.

Unit economics

WHAT	LENGTH	COST
One client clip	5 s	\$0.35
One client clip	10 s	\$0.70
The whole family, end to end	6 × 10 s	\$4.20

The **\$0.07** per second rate is confirmed by four independent billing measurements – at three, five and ten seconds. Everything else costs cents: seven of eight steps are free and sit before the paid one, so defects are caught on a still rather than on finished video.

It all comes down to the driving and the prompting

This is not a general observation — we proved it ourselves, one measurement at a time, and every time the lever turned out to be in one of those two places.

The prompt described a face — and erased the client's

We cut anthropometry out of the prompts: distance to the client moved from 0.44 to 0.36, and the blonde stopped turning brunette.

The prompt set the wardrobe — and put a skirt on a man

Wardrobe arrives from the text, not from the demo identity. Hence a template is bound to a gender as a whole, and the gate fires before generation.

The driving framing never reached the prompt — and the character drifted off frame

We started measuring the framing off the video and carrying it into the aesthetic prompt in words. Four templates of six landed in frame immediately.

A small face in the driving — and identity drifts

Measured across four drivings: 34–56 px against 96–139 px, and the results split exactly on that line.

None of these levers is inside the model. All four are in the material and in the text — where the work costs hours, not months.

Python + Kling Motion Control 6 templates • 2 identities • 4 drivings

Every number from a live measurement